




MONASH University

Information Technology

# FIT5202 (Volume I - Introduction)

Week 2a – Introduction to Parallel Databases

**algorithm** distributed systems **database**  
systems **computation** knowledge ma  
**design** e-business **model** data mining **int**  
distributed systems **database** software  
**computation** knowledge management **an**



# Chapter 1 Introduction


- 1.1 A Brief Overview - Parallel Databases and Grid Databases
- 1.2 Parallel Query Processing: Motivations
- 1.3 Parallel Query Processing: Objectives
- 1.4 Forms of Parallelism
- 1.5 Parallel Database Architectures
- 1.6 Grid Database Architecture
- 1.7 Structure of this Book
- 1.8 Summary
- 1.9 Bibliographical Notes
- 1.10 Exercises

## 1.3. Objectives (cont'd)

### • Parallel Obstacles

- Start-up and Consolidation costs,
- Interference and Communication, and
- Skew


Recall:



## 1.3. Objectives (cont'd)

### . Start-up and Consolidation

- Start up: initiation of multiple processes
- Consolidation: the cost for collecting results obtained from each processor by a host processor

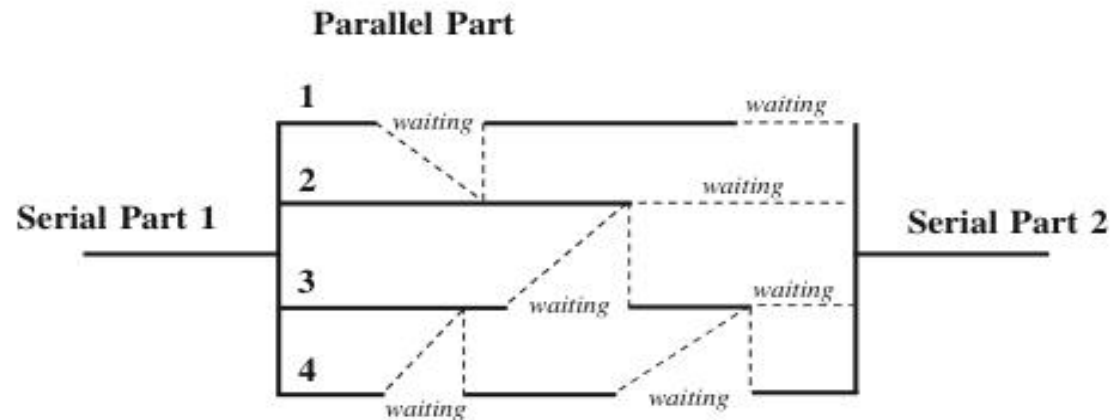


**Figure 1.3** Serial part vs. parallel part

## 1.3. Objectives (cont'd)

### . Interference and Communication

- Interference: competing to access shared resources (e.g., disk & memory)
- Communication: one process communicating with other processes, and often one has to wait for others to be ready for communication (i.e. waiting time).




**Figure 1.4** Waiting period

## . Exercise

- There is a job that will take 1 hour to complete, if this is done by 1 processor.
- The serial part of this job is 10%
- There are 4 processors to use in this job, but each processor will have an overhead of 20% due to waiting time, communication time, etc.
- What type of **speed up** do we get?

**There is a job that will take 1 hour to complete, if this is done by 1 processor. The serial part of this job is 10%. There are 4 processors to use in this job, but each processor will have an overhead of 20% due to waiting time, communication time, etc. What is the speed up? (5 Minutes)**

$$\text{Speed up} = \frac{\text{elapsed time on uniprocessor}}{\text{elapsed time on multiprocessors}}$$



**Figure 1.1** Speed up

**Solution:**

1 processor = 60min

Serial part = 10% = 6min

Parallel part = 54min

4 processors = 54min/4 = 13.5min

Overhead = 20%

Hence, parallel processing part = 13.5min + 20%overhead = 13.5min+2.7min = 16.2min

**Total time** = 6min (serial) + 16.2min (parallel) = 22.2min

Speed up = 60min / 22.2min = 2.7

Linear speedup should be 4


Speed up of 2.7 is Sub-Linear Speedup

## 1.3. Objectives (cont'd)

### • Skew

- Unevenness of workload
- Load balancing is one of the critical factors to achieve linear speed up

- ❑ **Data skew:** Uneven distribution of data in terms of size, allocated to processors
- ❑ **Processing skew:** Uneven processing time of the processors due to data skew.



**Figure 1.5** Balanced workload vs. unbalanced workload (skewed)



## 1.3. Objectives (cont'd)

### . Skew

- *Zipf* distribution model to model skew. Measured in terms of different sizes of fragments allocated to the processors

$$|R_i| = \frac{|R|}{i^\theta \times \sum_{j=1}^N \frac{1}{j^\theta}} \quad \text{where } 0 \leq \theta \leq 1 \quad (2.1)$$

- The symbol  $\theta$  denotes the degree of skewness, where  $\theta = 0$  indicates no skew, and  **$\theta = 1$  indicates highly skewed**
- $|R|$  is number of records in the table,  **$|R_i|$  is number of records in processor  $i$** , and  $N$  is number of processor ( $j$  is a loop counter, starting from 1 to  $N$ )
- Example:  $|R|=100,000$  records,  $N=8$  processors




Figure 2.1 Uniform distribution (no skew)





Figure 2.2 Highly skewed distribution

## 1.3. Objectives (cont'd)

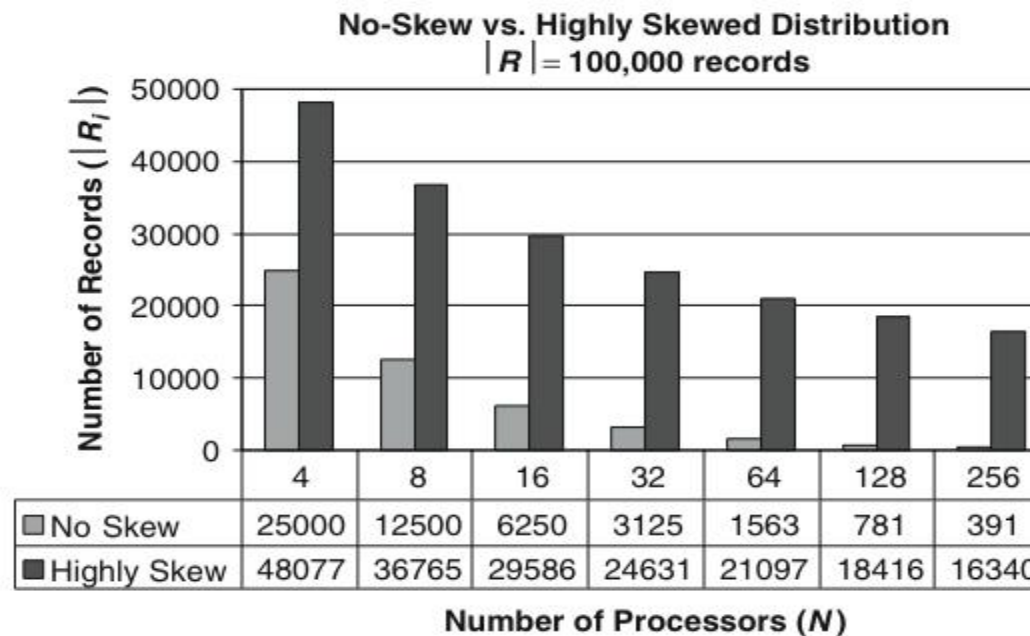
- No skew vs. highly skewed



**Figure 2.3** Comparison between highly skewed, less skewed, and no-skew distributions

## 1.3. Objectives (cont'd)

- No skew vs. highly skewed



**Figure 2.4** Comparison between the heaviest loaded processors using no-skew and highly skewed distributions

## 1.3. Objectives (cont'd)

- No skew vs. highly skewed

**Table 2.2** Divisors (with vs. without skew)

<i>N</i>	4	8	16	32	64	128	256
<b>Divisor without skew</b>	4	8	16	32	64	128	256
<b>Divisor with skew</b>	2.08	2.72	3.38	4.06	4.74	5.43	6.12

## . Exercise 7

- There 100,000 records in the table to be distributed to 32 processors. Assuming that the skewness degree is high ( $\theta = 1$ ), what is the estimated number of records in the heaviest processor?
- A. 48,000 records
- B. 29,000 records
- C. 24,000 records
- D. It is not possible to predict


## 1.4. Forms of Parallelism

- Forms of parallelism for database processing:
  - Interquery parallelism
  - Intraquery parallelism
  - Interoperation parallelism
  - Intraoperation parallelism
  - Mixed parallelism

## 1.4. Forms of Parallelism (cont'd)

### • Interquery Parallelism

- “Parallelism **among** queries”
- **Different queries** or transactions **are executed in parallel with one another**
- Main aim: scaling up transaction processing systems

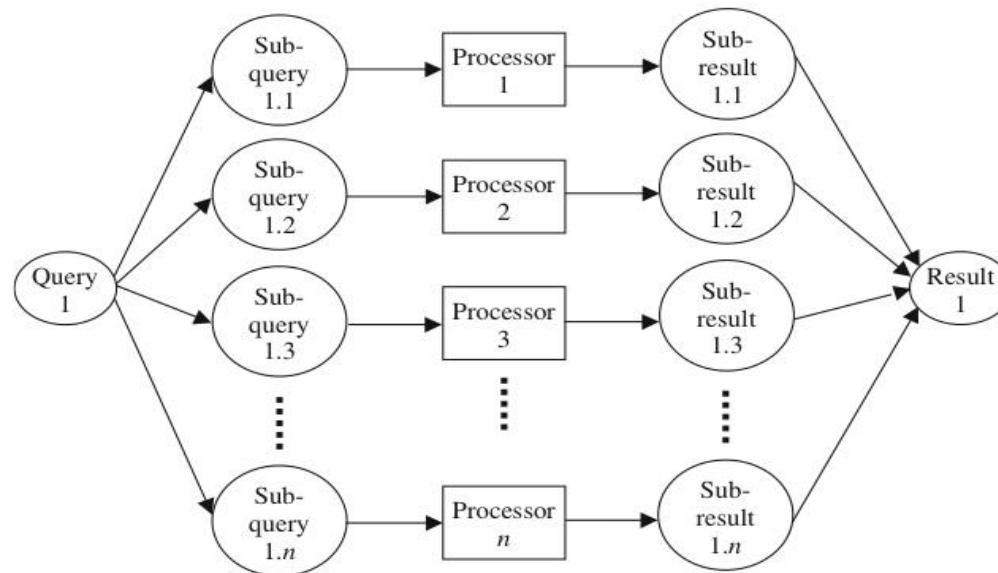


**Figure 1.6** Interquery parallelism

## 1.4. Forms of Parallelism (cont'd)

### • Intraquery Parallelism

- “Parallelism **within** a query”
- **Execution of a single query in parallel** on multiple processors and disks
- Main aim: speeding up long-running queries



**Figure 1.7** Intraquery parallelism



## 1.4. Forms of Parallelism (cont'd)

- Execution of a single query can be parallelized in two ways:
  - **Intraoperation parallelism**: Speeding up the processing of a query by **parallelizing the execution of each individual operation** (e.g. parallel sort, parallel search, etc)
  - **Interoperation parallelism**: Speeding up the processing of a query by **executing in parallel different operations in a query expression** (e.g. simultaneous sorting or searching)

## 1.4. Forms of Parallelism (cont'd)

### • Intraoperation Parallelism

- Parallelize execution of each individual operation
- “Partitioned parallelism”
- Parallelism due to the data being partitioned
- Since the number of records in a table can be large, the degree of parallelism is potentially enormous

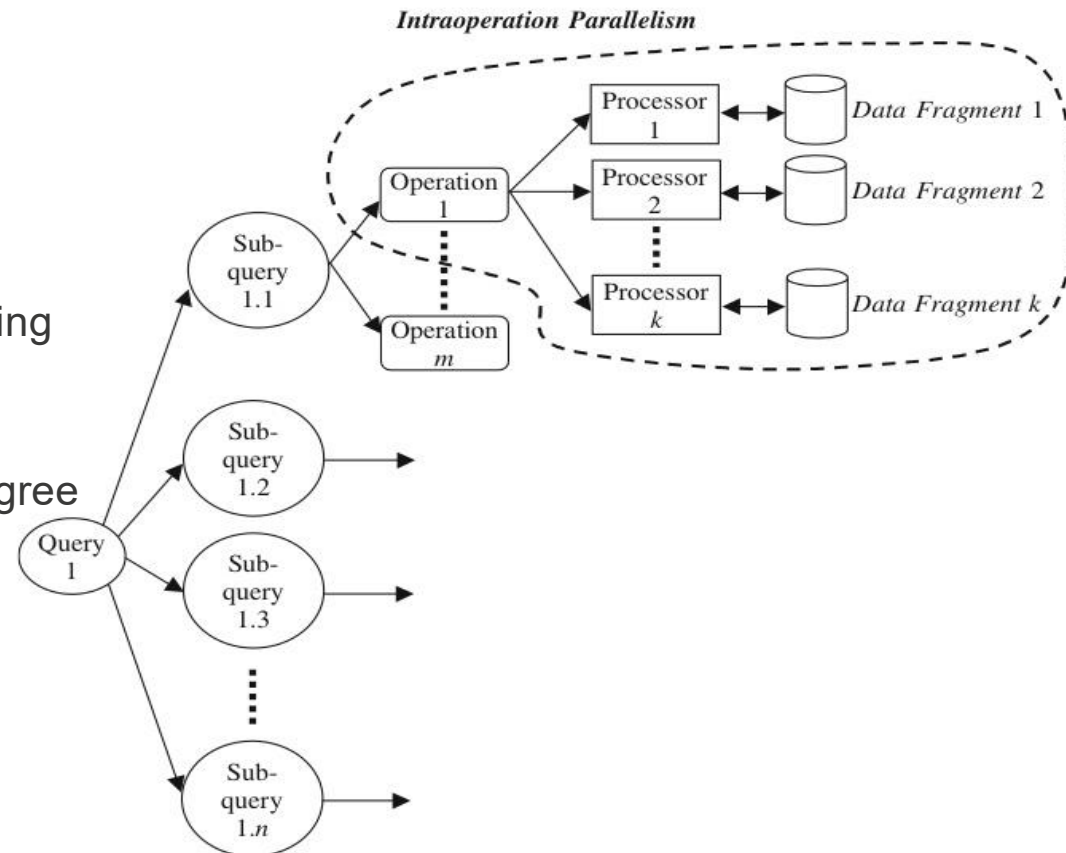


Figure 1.8 Intraoperation parallelism

## 1.4. Forms of Parallelism (cont'd)

- **Interoperation parallelism:** Parallelism created by concurrently executing different operations within the same query or transaction
  - Different operations are executed in parallel
    - Pipeline parallelism
    - Independent parallelism

## 1.4. Forms of Parallelism (cont'd)

### • Pipeline Parallelism

- Output record of one operation *A* are consumed by a second operation *B*, even before the first operation has produced the entire set of records in its output
- Multiple operations form some sort of assembly line to manufacture the query results
- Useful with a small number of processors, but does not scale up well

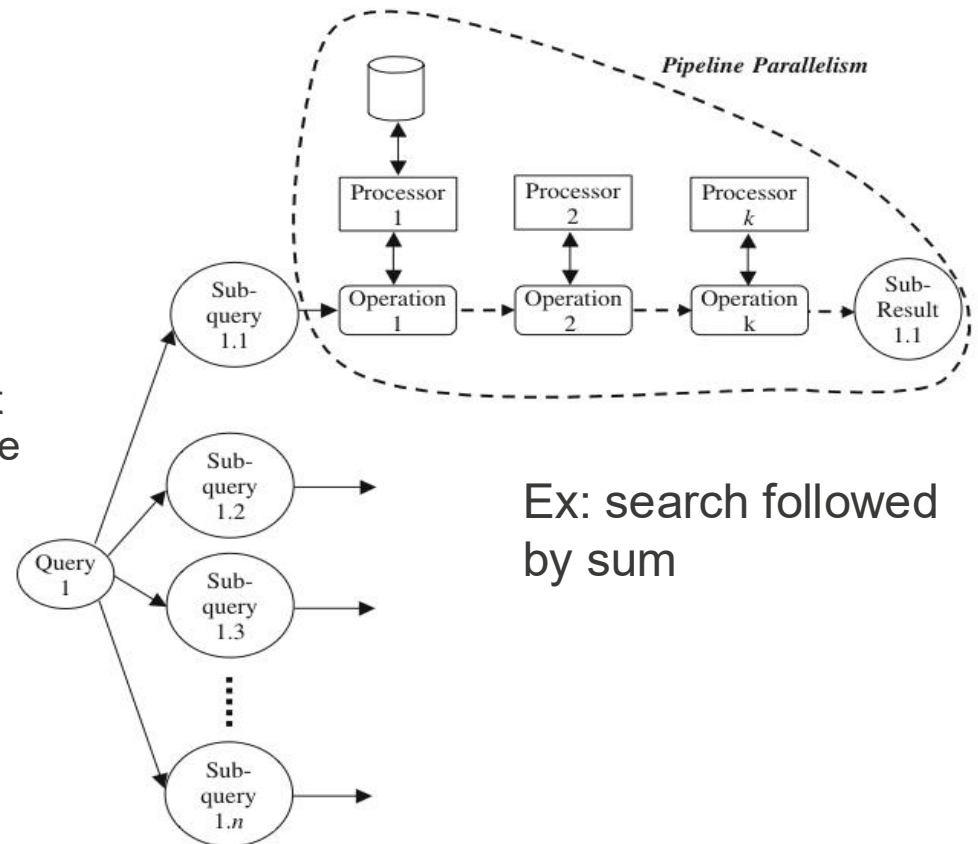


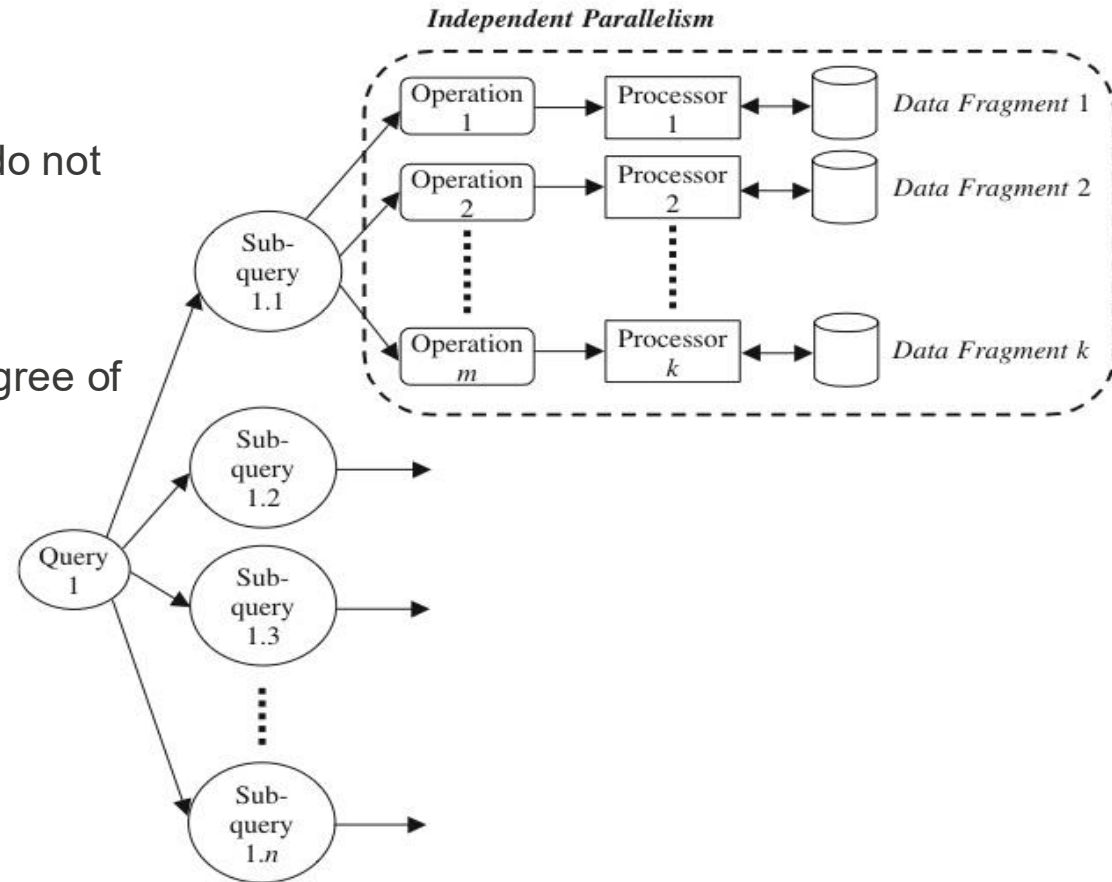
Figure 1.9 Pipeline parallelism

## 1.4. Forms of Parallelism (cont'd)

### Independent Parallelism

- Operations in a query that do not depend on one another are executed in parallel
- Does not provide a high degree of parallelism

Parallelism is limited by the number of operations available

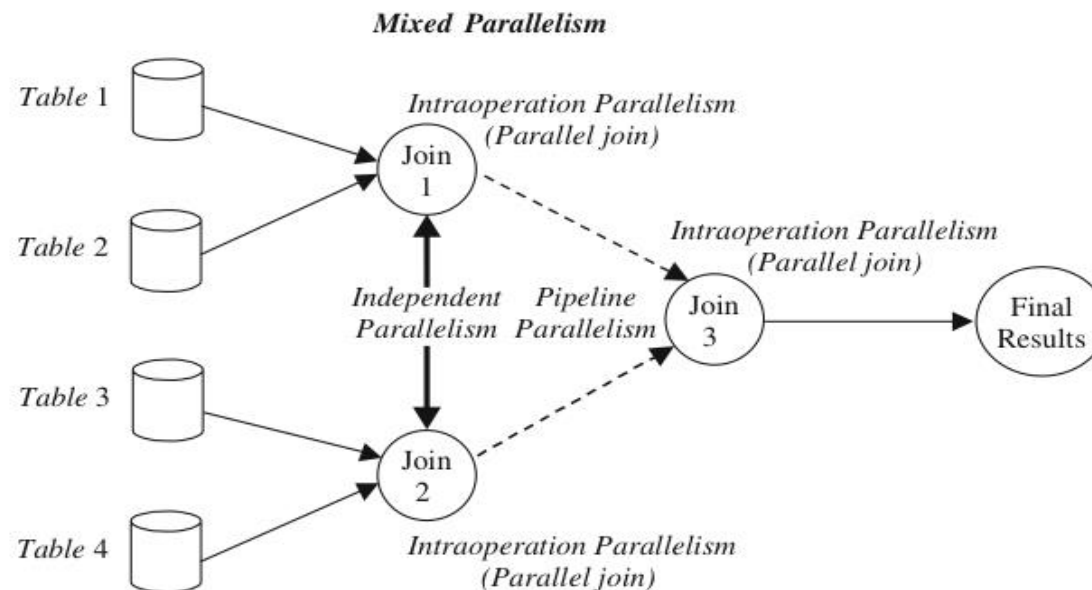


**Figure 1.10** Independent parallelism

## 1.4. Forms of Parallelism (cont'd)

### • Mixed Parallelism

- In practice, a mixture of all available parallelism forms is used.



**Figure 1.11** Mixed parallelism


## 1.5. Parallel Database Architectures

- Parallel computers are no longer a monopoly of supercomputers
- Parallel computers are available in many forms:
  - Shared-memory architecture
  - Shared-disk architecture
  - Shared-nothing architecture
  - Shared-something architecture

## 1.5. Parallel Database Architectures (cont'd)

### • Shared-Memory and Shared-Disk Architectures

- Shared-Memory: all processors share a common main memory and secondary memory
- Load balancing is relatively easy to achieve, but suffer from memory and bus contention
- Shared-Disk: all processors, each of which has its own local main memory, share the disks



**Figure 1.12** An SMP architecture


**Load balancing:** process of distributing a set of tasks over a set of computing units



## 1.5. Parallel Database Architectures (cont'd)

### • Shared-Nothing Architecture

- Each processor has its own local main memory and disks
- Load balancing becomes difficult (because data is placed locally in each processor, and each processor may have an unequal load)

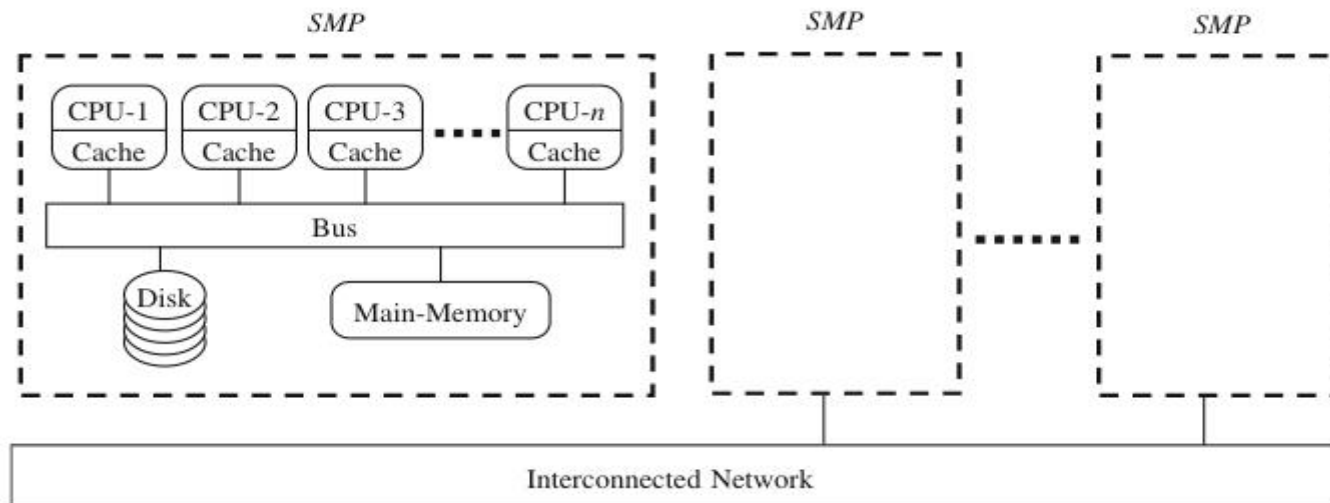


**Figure 1.13** A shared-nothing architecture

## 1.5. Parallel Database Architectures (cont'd)

### Shared-Something Architecture

- A mixture of shared-memory and shared-nothing architectures
- Each node is a shared-memory architecture connected to an interconnection network aka shared-nothing architecture

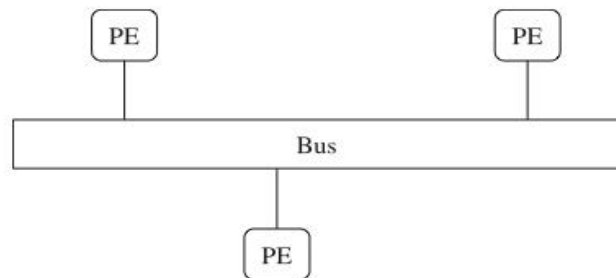


**Figure 1.14** Cluster of SMP architectures


## 1.5. Parallel Database Architectures (cont'd)

### • Interconnection Networks

- Bus, Mesh, Hypercube




**Figure 1.15** Bus interconnection network





**Figure 1.16** Mesh interconnection network

*2-dimensional*



*3-dimensional*



**Figure 1.17** Hypercube interconnection network

## 1.8. Summary

- **Why**, **What**, and **How** of parallel query processing:
  - Why is parallelism necessary in database processing?
  - What can be achieved by parallelism in database processing?
  - How parallelism performed in database processing?
  - What facilities of parallel computing can be used?